Applicability of SSM and UML for Designing a Search Application for the British Broadcasting Corporation (BBC)

Ross Fenning[1], Huseyin Dogan[2], and Keith Phalp[2]

[1] BBC, UK `Ross.Fenning@bbc.co.uk`
[2] Bournemouth University, UK {`hdogan,kphalp`}`@bournemouth.ac.uk`

Abstract. Whilst there are successful general web search engines such as Google that will find any piece of content, there is a perceived need for a specific search that makes better use of the internal knowledge the broadcasting industry (e.g. BBC) has about its own content. The British Broadcasting Corporation (BBC) is a public service broadcaster funded by the licence fee paid by United Kingdom households. This industry-based case study looks at the applicability of Soft Systems Methodology (SSM) and Unified Modelling Language (UML) to design a hypothetical, high-level view of a search application that receives web content from a variety of BBC content production systems and makes every item then searchable by a BBC website visitor using the search feature. The developers of such search applications can benefit from this specific industry-based case study that contextualised the problem space using SSM and developed UML models to solve the problem.

1 The Problem

1.1 Background

The BBC has been publishing content on the World Wide Web since the mid 1990s and since then the amount and the diversity has increased exponentially. Large websites – or indeed the web as a whole – would not have been usable nor useful without the rise in quality of web search engines.

A large challenge for any web search application is to provide a common interface and set of user interactions that can equally index, search and link to a diverse range of types of information – be it in the form of text, images, video or games. A more recent challenge has been to achieve this in a near-real time way to catch up with the rapid rate at which content is added to the web (particulary from microblogging websites such as Twitter).

Whilst there are successful general web search engines such a Google that will find any piece of content, there is a perceived need for a BBC-specific search that makes better use of the internal knowledge the BBC has about its own content.

1.2 Design Goals

Whilst the BBC website currently has a functional search feature already, this paper assumes building a new search application from the ground up so as to give full freedom to apply the analysis and design techniques therein. In practice, there are engineering challenges involved in maintaining and building on top of existing systems. Such challenges and details of the existing infrastructure are out of scope for this paper.

The purpose of this paper is to explore how contextualisation through SSM and then using UML to form a high-level design might benefit building a hypothetical, new search application for the BBC. The hope is that insights from fresh analysis and design might form suitable proposals for potential improvements and development within the existing application. These approaches might also lead to further work to evaluate the current application by highlighting where the BBC Search application differs from an "ideal" model (if it can be described as such) derived from such a redesign with a contemporary view of the problem space.

Ultimately, any organisation like the BBC is unlikely to replace a large application atomically, but is likely to migrate over time to its "ideal" form with smaller, iterative improvements. It is proposed that there is some value in designing what that "ideal" form might be so as to provide proposals for those improvements.

1.3 Problem Space Contextualisation through Soft Systems

The target audience for the BBC is effectively the entire population of the UK and amongst those that do make use of BBC services, there is much diversity of needs, preferences and technical ability. It is clear it is no small task to design a search-based discovery mechanism of millions of diverse pieces of content aimed at millions of diverse people.

Using *Soft Systems Methodology* (SSM) [1], we can stand back from an ontological approach of defining what the search system *is* or *comprises* and instead take an *epistemological* view of search as a system. With this view, we could consider a system that holistically transforms members of the public's desires to find online content into the consumption of that content – whether those desires are *precise* (e.g. they want an exact article known by headline they saw earlier or a particular programme they missed on television) or those desires are *fuzzy* (e.g. news about a certain topic, any comedy programme, learning materials about the Industrial Revolution).

Checkland and Scoles[2] decribed a *Rich Picture* approach for representing a problem situation early in SSM approaches. Given the size and complexity of the search system as a whole, a useful initial step is to create such an informal representation of what is known about the problem. Figure 1 shows what the authors know of the audience, search and most BBC online content areas. Note that not all areas are covered and a strong emphasis is placed on TV catch-up (e.g. via the iPlayer product). Radio catch-up is not mentioned as it shares a lot

of similarity with television in terms of use and any differences are out of scope for this design.

This rich picture was created using domain knowledge dervied from the existing search system and technical knowledge of the BBC websites. If the initial designs from this paper prove promising in practice, the authors would recommend repeating the exercise with a broader range of stakeholders and domain experts from respective subsystems. A rich picture does not have to serve as authoritative snapshot of the problem space, but can instead been seen as a collaborative exercise between stakeholders. Such contextualisation is likely worth iterating over time as the industry changes, e.g. the recent shift from separate mobile websites to *Responsive Design*[3] might change the understanding around mobile devices.

Dogan and Henshaw[4] showed how a "soft" systems approach called Interactive Management can be adopted to capture the requirements and contextualise the problem space. This involved the process of transitioning from the soft systems results to a formal model (e.g. UML). This transition was enabled by dividing the actors in the rich pictures into meta-level and direct users of the system. Although this division was subjective and depended on the interpretation and analysis of the rich pictures to provide a structure for the use case model, the rich pictures themselves were created through interactions with subject matter experts. The soft systems, and hence the Interactive Management results, provided the baseline information to derive a formal model including UML use case, sequence and domain models. The transitioning from "soft" to "hard" systems can be set within the State of the Art including the requirements analysis and modelling as used in SSM, UML and Business Process Modelling.

The overall design objective of this paper will be to create an initial proposal for a search application to drive the missing components within the holistic system depicted in Fig. 1. Some subsystems already exist, e.g. for journalists to write news articles and publish them on the BBC News website, but for the purposes of this design exercise, we will assume no existing application to drive a search-based discovery of those websites.

The design will look to integrate with existing subsystems where possible rather than attempt to replicate work already done. For example, journalists will prefer that a search application can integrate with the system into which they are publishing their articles instead of being required to publish their articles into two systems.

2 Design

2.1 Use Cases

From the rich picture in Fig. 1, we extracted the activities that are clearly within the remit of a search application. For example, the ability for editorial staff to manage the content of the search indexes sounds like a feature the search application would provide. Conversely, television actors and other contributors to
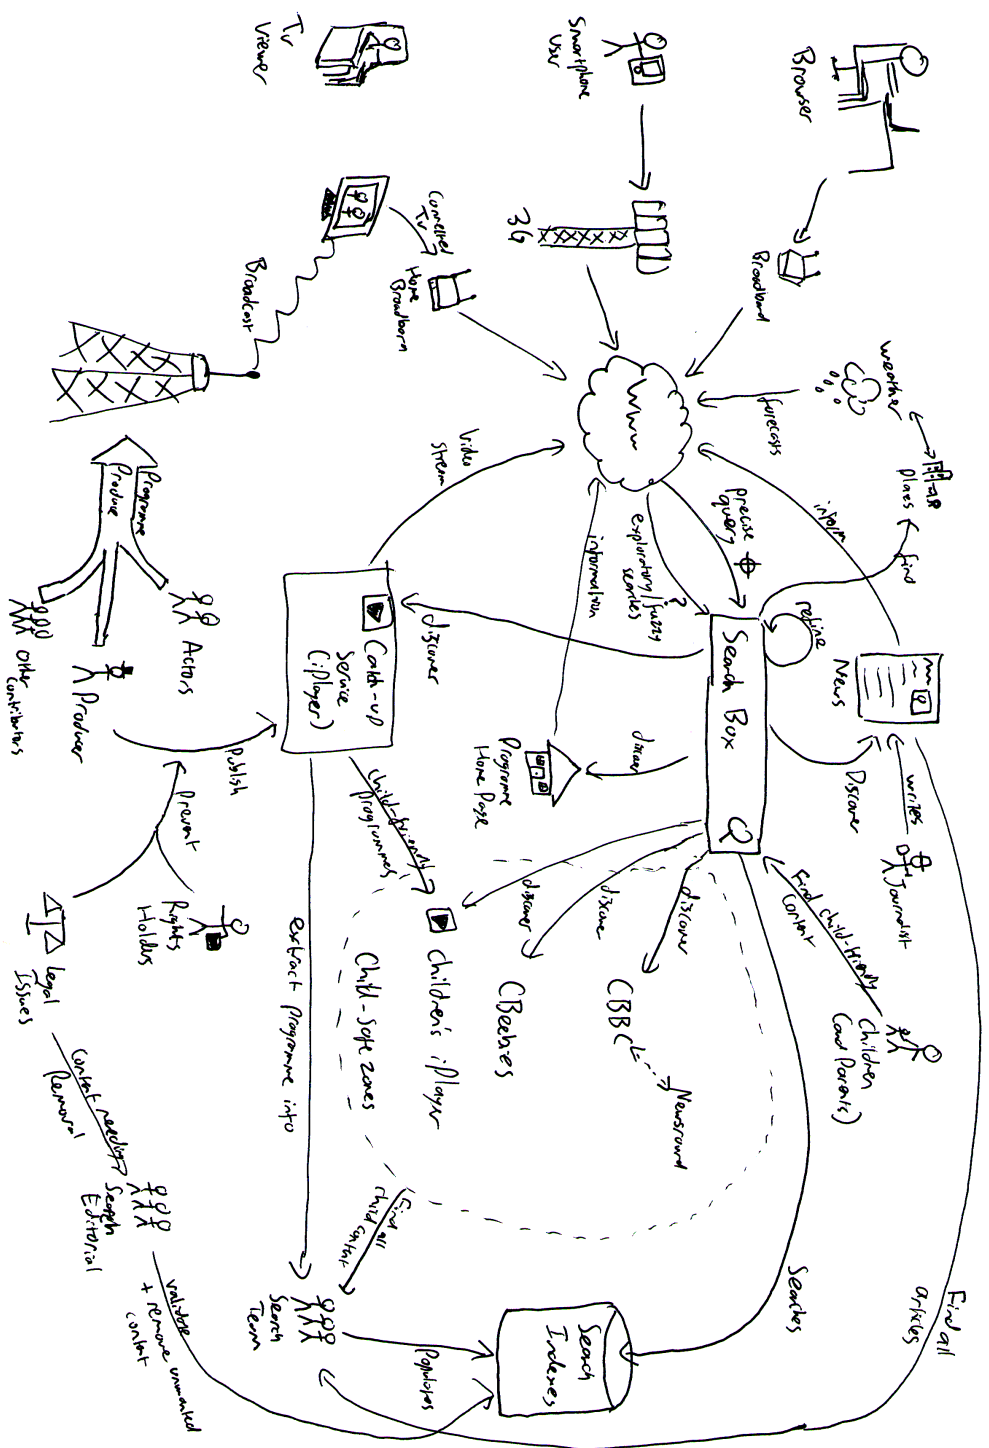
Fig. 1. Rich Picture of the BBC search system

a programme are likely to interact only within a television production subsystem with a producer or content editor being responsible for publishing information about the final production's broadcast and availability for streaming online.

In other words, we can say that editorial staff and content editors might be seen as direct users of the system, with actors, presenters, etc. seen as meta-level users. This is justified by noting that journalists and editorial staff will have their content exactly as typed appear in search results, but production staff behind television and radio programming contribute only indirectly to the search system. They produce the content that users will later wish to see, but they have little contribution to the discovery of that content later on.

This follows from the Interactive Management approach from Dogan and Henshaw[4] and is still largely subjective. For example, if a journalist publishes to a content management system, with which a search system then integrates without their knowledge, can they still be considered to be a direct user of the search system? In this design we argue that such syndication into the search system means their actions (e.g. to publish, remove, update articles) will have direct effect on the search and discovery of those articles (e.g. they may become searchable, cease being searchable or start being searchable under new criteria) and thus they are direct users of the search system. It is recognised however, that a more collaborative rich picture drawn up with a wider set of stakeholders might lead to a different opinion.

This is not to say that we can simply cross off certain elements from our depiction of the problem because they do not directly interact with the subsystem being designed. The *systems thinking* approach advocated by Checkland[5] encourages us to consider the irreducible properties of each system at each level of abstraction. Thus we need to consider not only a search application subsystem that solves specific problems for its immediate users, but also an application that contributes to the desirable, emergent properties of the BBC service as a whole.

In the specific example of television programming, we need to maintain systems thinking throughout the design process to ensure that we create a search application that both meets the needs of the public using the application to search for programmes and forms part of a television production and delivery system that itself meets the needs of the television-watching public.

Thus a suitable design strategy is to apply systems design to the search application in isolation – as a *hard problem* – but then to use the wider system to inform, shape and evaluate that design.

Having extraced the direct users only of the search system, further requirements-gathering and business analysis can define their respective use cases of the system. An analysis based on domain knowledge from the existing search system leads to the use case UML diagram shown in Fig. 2.

This illustrates only a subset of the expected behaviours for a full BBC search application, but touches on some of the diversity of the potential uses. For the purposes of our initial design, we can next look into defining the system behaviour for some of these use cases. It should be noted that these are illustrative of the breadth of use cases, but deeper business analysis and creation of

corresponding use cases (perhaps following the style of Cockburn[6]) should be performed as part of a more in-depth study before the high-level design can be truly considered complete and valid.

2.2 System Behaviour

The breadth of content the search system will need to index and retrieve sits across multiple systems within the BBC (e.g. programme metadata lives in a distinct location to news articles). The diversity of systems, formats and metadata involved suggested approaching the more functional design as a classic *Enterprise Integration* problem[7].

Figure 3 shows how a TV producer indirectly interacts with search by providing information that ultimately ends up in the search indexes and Fig. 4 shows a sequence diagram defining a user interacting with the search system.

The key design decision in Fig. 3 is the use of asynchronous messages only. A non-blocking set of interactions such as publish-subscribe[8] is a good way to decouple systems that produce and store progamme information from the search application systems. If the systems surrounding the programmes database can be built a *channel adapter*[8] to integrate it to a messaging system, then the search indexes can receive changes to information without TV producers, journalists, content editors, etc. even being aware this is happening.

Note in Fig. 4 that while the search indexes will contain representations of content from several source systems, the intention shown is that richer information about the domain model will not be held in the search indexes. This goes along with the principle of separation of concerns[9] in that the search index component can focus on optimising its data structures around retrieval. This kind of modularity also allows source systems maintained by other teams to take the responsibility of the accuracy of the information, which fits in with the wider holistic view of the system: the volume of information involved requires that separate teams are responsible for the accuracy of their data.

This can be seen as similar to the *Lazy Load* pattern [10] in that the search indexes will only return stub objects that are capable of retrieving the fuller information to need. This, however, could lead to a lot of calls to different service applications per page of results. This can be done in parallel, but the fact still remains that the user has to wait for all this to assemble before seeing even one result.

One solution to this in modern web application design is to push the lazy loading into the web browser using AJAX[11]. Such a solution is depicted in Fig. 5. Whilst this still requires just as many calls to backend services – perhaps even more complexity as the calls are going through more layers – it can give a user experience that appears more responsive.

Preparing a search results page with minimal information that is then augmented asynchronously is a user experience technique that gives the illusion of lower latency; the additional information can update the page during the user's reaction time in the best case.
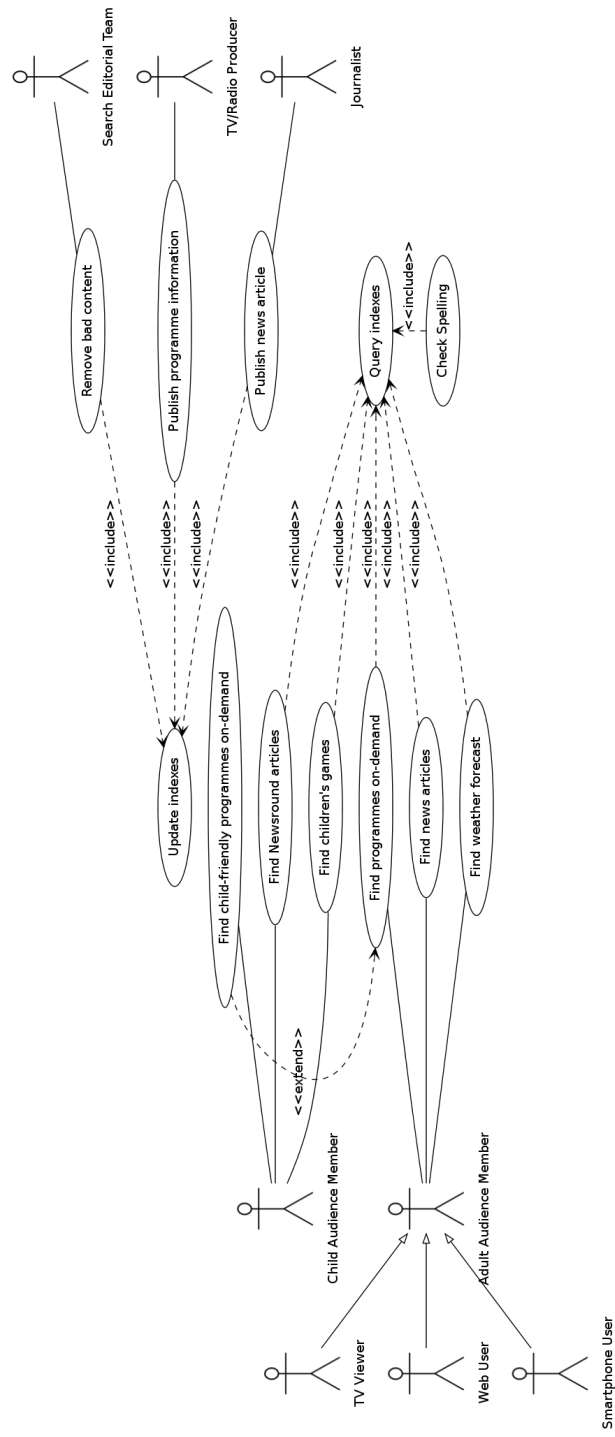
Fig. 2. Use case diagram for BBC Search application
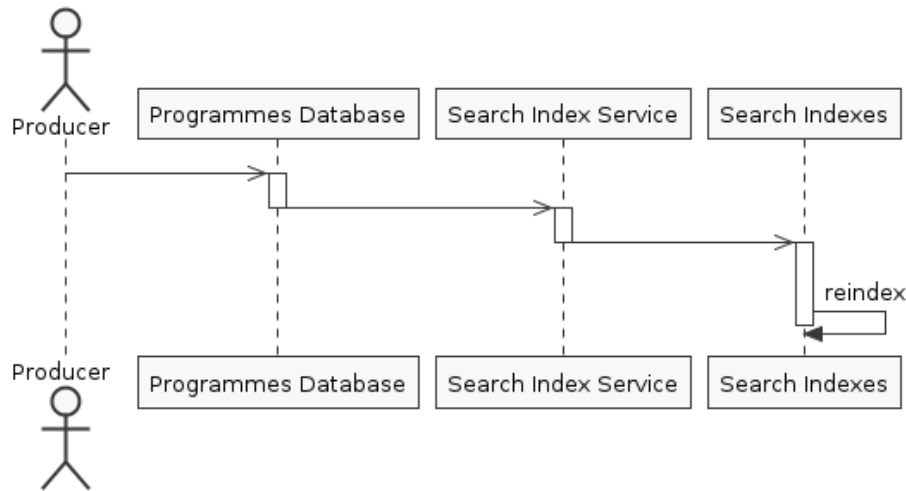
Fig. 3. Sequence diagram showing publication of programme information to the search indexes

2.3 Domain Model

It has been expressed already that a BBC-wide search application would need to index, retrieve and display a diverse set of information. Parts of the application might well want to incorporate a domain model[10] so as to understand how to perform each of these actions against each possible item the search application could return as a search result.

Domains are the distinct subject matters present in any system representing large, reusable components and are depicted using a domain model which shows an organisation of UML packages and their dependencies[12]. The use case diagram and domain model developed provides a baseline model for a future search application system.

A maximal domain model is shown in Fig. 6 that attempts to capture a good proportion of the content and concepts the BBC has been making efforts to model over several years. This model is an aggregate of individual ontologies developed for specific purposes, but given the search application has to provide the discovery for the full set of this information, it is not unreasonable that a search application would have a domain model that covers the totality.

Note that complexity of information for programmes alone[13]. A programme to a member of the public could actually refer to an exact episode or indeed the *brand*, i.e. the title of the programme in general. An example of a brand would be *Doctor Who*, [3] which itself contains multiple series, which in turn contain

---
[3] *Doctor Who* is a popular, long-running science-fiction television programme produced by the BBC and is frequently sought by users of the BBC iPlayer television catch-up service after an episode has aired.
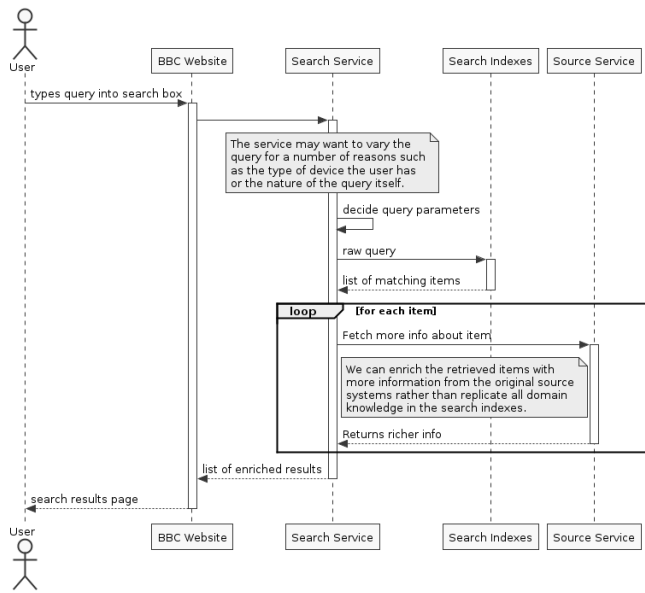
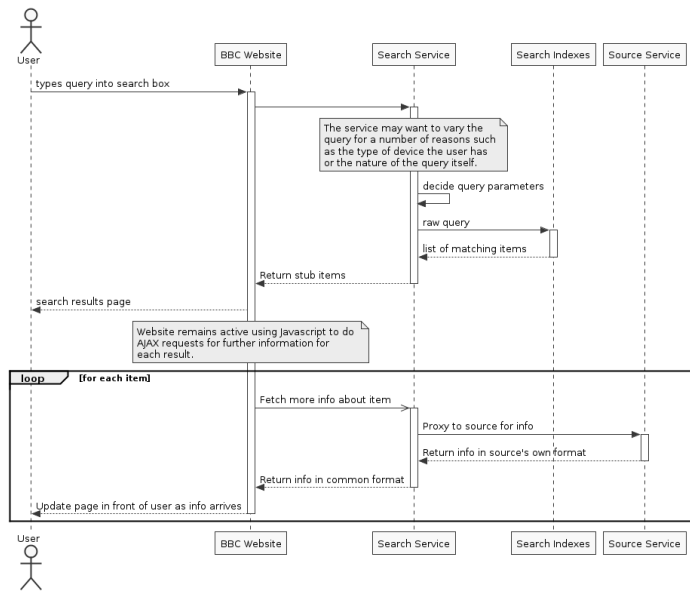**Fig. 4.** Sequence diagram showing a user interacting with search

Fig. 5. Sequence diagram showing rich information lazy-loaded via AJAX

a collection of episodes. Note that the brand itself has episodes as immediate children that do not live under a series, e.g. Christmas specials. It is also the case that certain things do not have brands or series, e.g. a film is modelled as a one-off episode.

This leads to some difficult questions for a search application. If a user searches for the text "doctor who", are they expecting a link to the latest episode to watch on iPlayer, information about the next episode – such as when it is due to be broadcast – or a link to the overall home page for the entire Doctor Who brand?

The model also skims the surface of the sport ontology[14] created before the 2012 Olympic games, which aims to model the whole domain of sporting personalities, events and competitions (and more). This might be too fine-grained for the domain model used within the search application, but it is likely that people will want to search for competitions like "World Cup" or sporting disciplines such as "football". A search application that understands these concepts as entities in their own right may well be able to direct users at a curated, dedicated "home page" thereof alongside simply matching articles and other works that contain those terms.

2.4 Discussion and Analysis

The problem of a BBC-wide search application was predicted – and has certainly shown itself – to be a very large-scale problem, the full extent of which cannot be covered in this paper. Thus it is reasonable to conclude that the problem is not yet solved at every level, but the high-level solutions are certainly promising.

The use case diagram is likely not complete and would need a significant amount of user research and requirements-gathering to collect all the possible use cases of the search system. It is likely that the variations of the use cases are so numerous that different diagrams exploring different combinations of use cases should be made to replace the single one given. For instance, very little has been touched on around users seeking educational and informative material such as *Bitesize* [4] or any other learning resources.

The sequence diagram for the producer (or any other content creator) indirectly getting their content into the search system demonstrates the need for asynchronous messaging (and publish-subscribe), but the nuances of such a process are not comprehensively shown in this format. A better illustration for such a message-based integration system might be derived from the illustrated patterns created by Hohpe and Woolf[8].

The behaviour of the query web application in Fig. 4 is defined with a hard line taken on keeping the information stored in the indexes to a minimum. At a basic level, this is not unlike a *content enricher* pattern [8] whereby the search application business layer enriches the stubs in the indexes with information the

---

[4] *Bitesize* is the name given to the BBC's free web-based study materials for school children aged between 5 and 16, covering varying curricula for England, Wales and Scotland.

Fig. 6. Domain model for content items and other things pertinent to BBC content

source system (the index) simply does not have. At a more purist extreme, this could follow a *claim check* pattern where the index returns only globally-unique identifiers for the matching items and makes no attempt to present any other knowledge about them (and thus avoids any synchronisation issues if an item is updated at source).

The latter pattern fully decouples the search indexes from any deeper information – leaving that responsibility to appropriate source system – but means that the raw, stub results are near-useless to an end user. This makes the AJAX-based alternative approach in Fig. 5 less of a desirable option (would a user really be presented a list of URIs while some Javascript code replaces them one by one with actual information?). Given the AJAX-driven behaviour still appears desirable in terms of responsiveness, it would seem that some trade-off would need to happen in terms of what the indexes store as additional information.

These decisions relate to how the system would implement the data model in Fig. 6. The ontology described is an aggregation of several efforts by Raimond[13], Rayfield [14] and others – along with some additional work to join them together – to represent the wealth of information with which the BBC deals. It could be argued both that a search application that *understands* this diversity of content must reflect it in its domain model but that a search application that is *coupled* to the individual ontologies in this way is brittle with regard to changes therein. Duplication of business models across different applications would only harm maintainability.

Thus the domain model given in Fig. 6 should only serve as a communication artefact that leads to further development of a domain model more suitable to the needs of the search application without the burden of maintaining more than is necessary. The use of *subtype polymorphism*[15] is likely to be key to ensure the search domain model contains only the APIs it needs. For example, does the search application have a need to differentiate between a *NewsItem* and a *BlogPost* for the purposes of displaying the search result's title? If the CreativeWork top-level class has a title property, then the domain model so far to enable that one behaviour needs only one class!

The application could go further and interact via a single SearchResult *facade*[10] whose instances provide appropriate responses to canonical hooks such as *getDisplayTitle()* and *getDestinationUrl()* via polymorphic *composition* with different target classes from the fuller domain model. Given that the search indexes are likely to take in content from any number of future systems, there is an appeal to taking a schemaless approach[16] and allow the search application's view layer to display different kinds of results differently via *duck typing* [17].

A dynamic, schema-free domain model used within the search application – noting that, of course, schemaless truly means there is an implicit schema in the logic that creates these dynamic objects from source systems[16] – could prove to aid the search system's need to model different kinds of data in the same index. For example, our designers might express a desire to put small summaries of weather forecasts within search results that contain places. In the static sense, we could say that anything of *Place* type is displayed with such a feature. With

a dynamic, duck-typed model, we could ask "does this item have weather?" in place of "is this a place?", thus allowing for the future design that extends forecasts to football matches (since football matches are *Event*s that occur in a *Place*, it is reasonable for a weather property to be set thereon).

Overall, we have some promising, high-level models from which to start making such more fine-grained decisions about the search application. The use cases are likely sufficient for early iterations or a *Minimal Viable Product*[18] and the behaviours capture the overall needs of the system. There is still much scope for returning to our systems thinking and Soft Systems Methodology to monitor the general model for Checkland's "3 Es" (efficacy, efficiency and effectiveness)[2], which is only briefly touched on in the authors' attempts to relate the data model to the user's interaction with the system.

3   Evaluation

The high-level nature of the modelling achieved has made it difficult to meet any objectives over performance, latency or robustness of software components involved. A BBC search application needs to index news as it is published, serve millions of unique visits every week and minimise potential for losing information. These non-functional requirements are a challenge in their own right and have not been met with the UML modelling presented.

This is not to say that UML is not capable as a tool for presenting architecture around performance and resilience. However, it is perhaps more appropriate to model and present some of these aspects through component and deployment diagrams. Behaviour diagrams such as sequence and activity do not suit well to showing timing or performance, although a series of sequence diagrams could show how behaviour changes in parts of the system to tolerate failure of other components (e.g. one component could be modelled to return from an internal cache if a collaborator is returning temporary error status codes).

Even the models that are presented in this paper do not paint the full picture, but it could be argued that it is not their purpose to do so. As stated at the end of Sect. 2.4, the use case diagram provides a starting point that might be sufficient for an early iteration of a project in an *Agile* methodology[19].

An Agile approach to developing the search application could distill the use cases even further and shape the requirements to the rest thereof later in the development process based on feedback and reacting to changes. A similar approach would allow us to start with the smaller domain model also discussed in Sect. 2.4 and allow it to grow to the necessary size to need – i.e. we can defer the decision of "how big should the domain model be?" until we are at a point where we have more information to answer such a question.

Thus even if the use case and domain models are not as comprehensive or as honed as they need to be to build an entire application, they serve their purpose adequately to communicate the first iterations of development or – especially in the case of the domain model – information pertaining to the whole organisation,

even if only small parts of it are they modelled directly in the application being built.

The sequence diagrams seem to provoke more debate of the merits of returning rich information in a single response versus an AJAX-driven approach (or some trade-off in between). Again, the authors would emphasise the communication aspect of UML modelling and argue that encouraging such debate is a successful application of UML, not a failure because a decision has not been made between two designs at this early stage.

Some of this hints at a drawback of UML modelling being that it encourages a lot of design decision at an early stage of a project – a stage at which we arguably have the least information[20]. However, there are plenty of efforts in spite of this that promote modelling and UML being used compatibly within an Agile process. The use of other modelling techniques in an Agile setting such as Rational Unified Process (RUP) and Agile Unified Process (AUP) have been suggested as suitable in an Agile project[21]. It may have been more suitable to incorporate a more diverse range of modelling techniques in the design and analysis presented so as to find the true strengths of each respective approach.

In conclusion, the UML modelling presented communicates some promising approaches to the BBC search problem, but it is far from sufficient in its own right. Designing to the level of detail required would result in a rigid development plan that made unverified assumptions, but an iterative approach to modelling that starts at a high-level and updates as development progresses could be used successfully in such a project.

A post-analysis study is required to evaluate the SSM and UML models. The models need to be applied to further projects for validation and verification purposes. In addition, the maturity and evolution of the artefacts need to be considered e.g. adding, deleting, or modifying, if the boundary and context changes.

References

1. Checkland, P., Poulter, J.: Learning for action: a short definitive account of soft systems methodology and its use for practitioner, teachers, and students. Volume 26. Wiley Chichester (2006)
2. Checkland, P., Scholes, J.: Soft systems methodology in action. Volume 7. Wiley Chichester (1990)
3. Marcotte, E.: Responsive web design. A list apart **306** (2010)
4. Dogan, H., Henshaw, M.: Transition from soft systems to an enterprise knowledge management architecture. In: International Conference on Contemporary Ergonomics and Human Factors 13-15 April 2010 Keele University, UK. (2010)
5. Checkland, P.: Systems thinking, systems practice: includes a 30-year retrospective. John Wiley & Sons (1999)
6. Cockburn, A.: Writing effective use cases. Pearson Education (2001)
7. Brosey, W.D., Neal, R.E., Marks, D.F.: Grand challenges of enterprise integration. In: Emerging Technologies and Factory Automation, 2001. Proceedings. 2001 8th IEEE International Conference on. Volume 2., IEEE (2001) 221–227

8. Hohpe, G., Woolf, B.: Enterprise integration patterns: Designing, building, and deploying messaging solutions. Addison-Wesley Professional (2004)
9. Dijkstra, E.W.: On the role of scientific thought. In: Selected Writings on Computing: A Personal Perspective. Springer (1982) 60–66
10. Fowler, M.: Patterns of enterprise application architecture. Addison-Wesley Longman Publishing Co., Inc. (2002)
11. Garrett, J.J.: Ajax: A new approach to web applications. `http://www.adaptivepath.com/ideas/ajax-new-approach-web-applications` (2005)
12. Dickerson, C.E., Mavris, D.N.: Architecture and principles of systems engineering. Complex and Enterprise Systems Engineering Series. Auerbach Publications (2009)
13. Raimond, Y., Sinclair, P., Humfrey, N., Smethurst, M.: Bbc programmes ontology. `http://purl.org/ontology/po/2009-09-07.shtml` (2009)
14. Rayfield, J., Wilton, P., Oliver, S.: Bbc sport ontology. `http://www.bbc.co.uk/ontologies/sport/2011-02-17.shtml` (2011)
15. Booch, G., Maksimchuk, R.A., Engle, M.W., Conallen, J., Houston, K.A., Young, B.J.: Object-Oriented Analysis and Design with Applications. Addison-Wesley (2007)
16. Sadalage, P.J., Fowler, M.: NoSQL distilled: a brief guide to the emerging world of polyglot persistence. Addison-Wesley (2012)
17. Python Software Foundation: duck-typing: Glossary – python v2.7.5 documentation. `http://docs.python.org/2/glossary.html#term-duck-typing` (2013)
18. Junk, W.S.: The dynamic balance between cost, schedule, features, and quality in software development projects (2000)
19. Beck, K., Beedle, M., van Bennekum, A., Cockburn, A., Cunningham, W., Fowler, M., Grenning, J., Highsmith, J., Hunt, A., Jeffries, R., et al.: The agile manifesto. The agile alliance **200**(1) (2001)
20. Kelly, A.: Conway's law v. software architecture. `http://allankelly.blogspot.co.uk/2013/03/conway-law-v-software-architecture.html` (2013)
21. Ambler, S.: Agile modeling: effective practices for extreme programming and the unified process. Wiley.com (2002)